

Digital Systems L2

– Computer Systems

Fang Yuan

School of Electronics Science and Engineering
Nanjing University
Nanjing 210046

2022

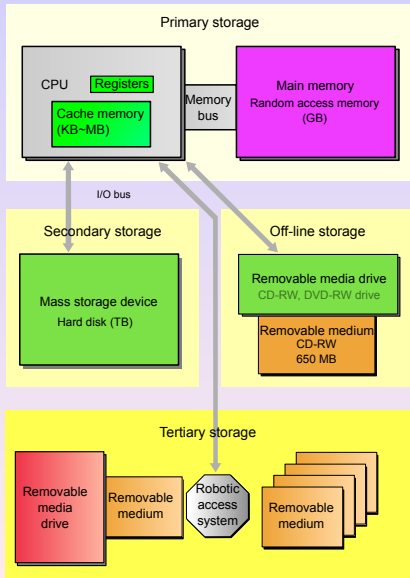


Outline

- 1 MEMORY AND FILE SYSTEMS
 - Memory Architecture
 - Cache
 - Virtual Memory
 - Filesystem



Types of Memory




- Primary storage vs. secondary storage.
- Volatile vs. Non-volatile.
- Static RAM vs. dynamic RAM.
- Types of ROM, from Masked-ROM through Flash ROM.
- Physical memory vs. virtual memory.




Memory Hierarchy

Memory Hierarchy



SRAM



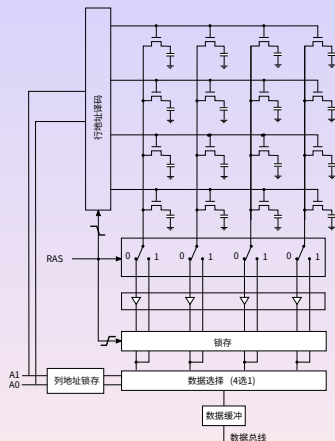
SRAM cell

READ BL and \overline{BL} set to HIGH, $WL_{HIGH} \rightarrow BL = Q$

WRITE set BL and \overline{BL} different level, $WL_{HIGH} \rightarrow Q = BL$



DRAM




4×4bit DRAM array

- higher density (lower cost per bit)
- capacitor needs recharge (refresh, more complex circuitry)



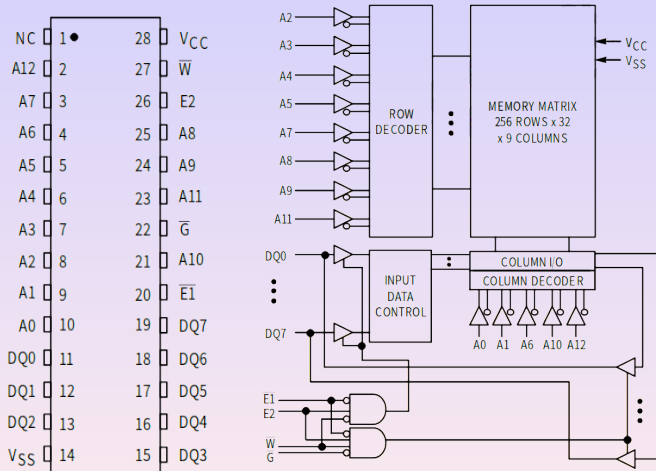
Memory array



Memory diagram

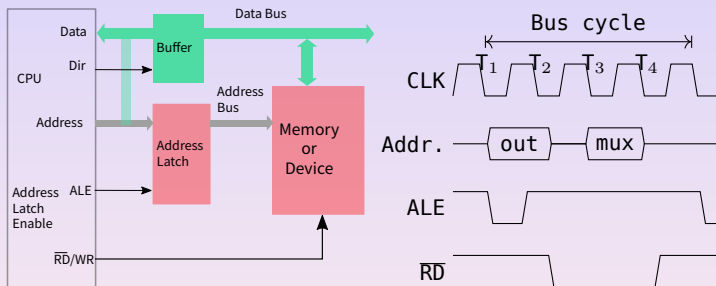


Static RAM




Bus Cycle

Bus Cycle is the cycle or time required to make a single read or write transaction between CPU and an external device (including memory).




Memory Timing

Read Cycle

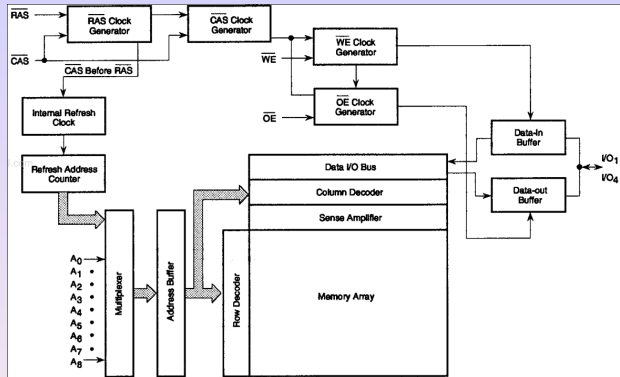



Memory Timing

Write Cycle



Dynamic RAM




DRAM (256K \times 4bit)




SDRAM(Synchronous DRAM)

- Synchronized with the system bus
- PC66-, PC100-, PC133- (66MHz, 100MHz, 133MHz)
- SIMM/DIMM(in-line memory modules) DDR(dual data rate)-SDRAM


t_{RCD}



CAS Latency



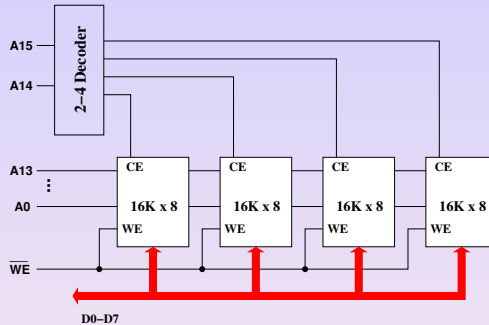
Words



Combining bits to words



Expansion

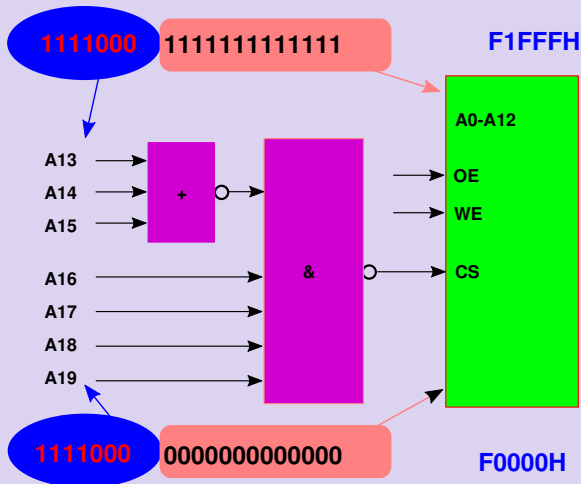


Memory Expansion




Full Decoding

8KB Memory decode



32-bit Bus

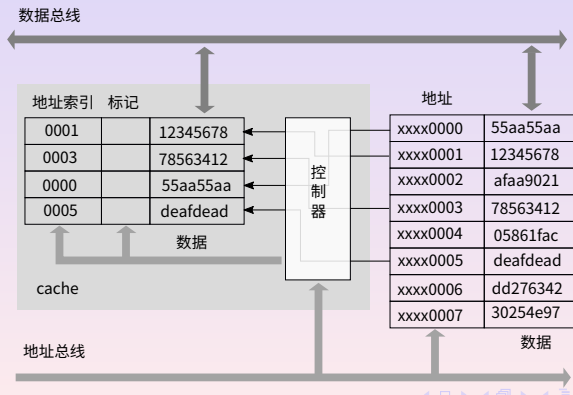


Memory system in 32bit computer



Why use Cache

- Memory and I/O need to be synchronized with CPU
- SRAM is faster than DRAM, but also expensive.



How Cache Works

Theory of Locality

Most programs trend to access memory within a limited range in a certain time:

- Spatial Locality
- Temporal Locality

Mapping DRAM to a small amount of SRAM (Cache) to accelerate access.


Available mappings:

- Fully associative cache
- Direct-mapped cache
- Set associative



Cache Map

Direct Map




Direct-mapped Cache



Cache Map

Set associative




2-way set associative



Consistency

- Reading

- **Look through**, always read from cache first




- Line fill



Consistency

- Reading
 - **Look aside**, read both from cache and main memory




- Line fill



Consistency

- Writing (updating main memory) when multi-processors or DMA present
 - **Write-through** at the same time, increase bus access




- **Invalidation**



Consistency

- Writing (updating main memory) when multi-processors or DMA present
 - **Write-back** at certain time (bus watching)



- **Invalidation**



What Happens?

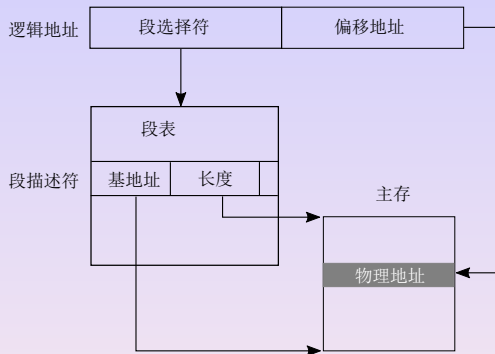
103ec: e3a02a02 MOV R2 , #8192

103f0: e5c23008 STRB R3 , [R2 , #8]

Read instructions and data



Segment Management

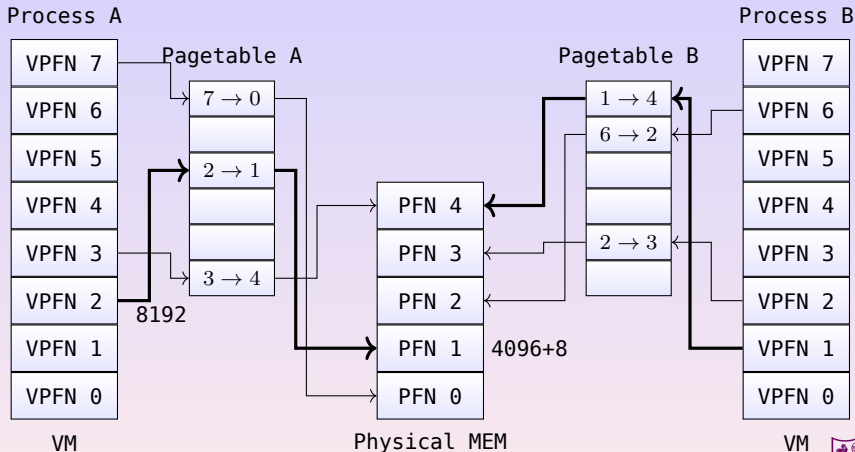


Address Conversion in Segment Unit

A logical (virtual) address consists of two parts: **segment selector** and **offset**.



Page Management



Paging Mode



X86(32bit) Segment Descriptor

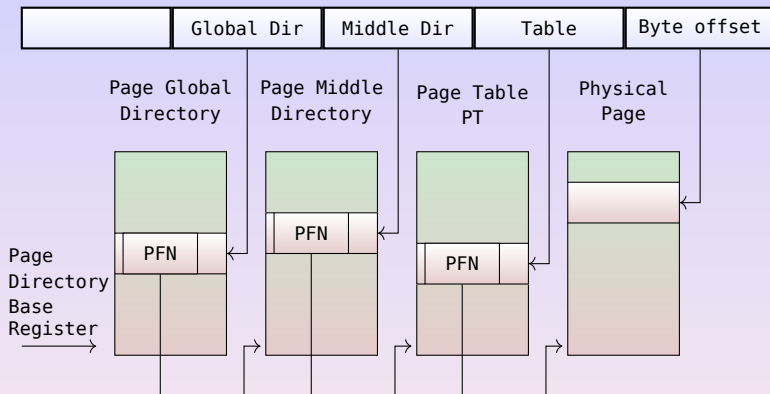
BASE ₃₁₋₂₄	G	D	0	AVL	LIMIT ₁₉₋₁₆	P	DPL	S	TYPE	A	BASE ₂₃₋₁₆
BASE ₁₅₋₀						LIMIT ₁₅₋₀					

Page Item

3	3	2	2	2	2	2	2	2	2	2	2	1	1	1	1	1	1	1	1	1	9	8	7	6	5	4	3	2	1	0			
Page Physical Base Address																				avl	G	P A T	D	A	P C D	P W T	U /	R /					



Multilevel Page Table



Three level of paging designing can be easily configured to 2 level mode by folding middle level.



TLB

Translation Look-aside Buffer (TLB) is a part of the chip's memory-management unit (MMU). TLB is used to reduce the time taken to access a user memory location.

Intel's Nehalem (2008, core i7/i5)

cache	4KB	2MB (pagesize)
L1	32KB+32KB	64 entries
L2	256KB (unified TLB)	512 entries/4way
L3	4–12MB	

example: hit – 1 clock cycle, missing – 30 clock cycles, missing rate – 2%.

average cycle: $1 \times 0.98 + 30 \times 0.02 \approx 1.6$ cycles



Page Swapping

Accessing an address that is not mapped to physical memory will trigger a page fault exception (interrupt).

- ① MMU finds a physical memory to be swapped (either drop or backup).
- ② The page that is swapped out is marked as “unmapped”.
- ③ Copy data.
- ④ Mark PFN in pagetable for this page.



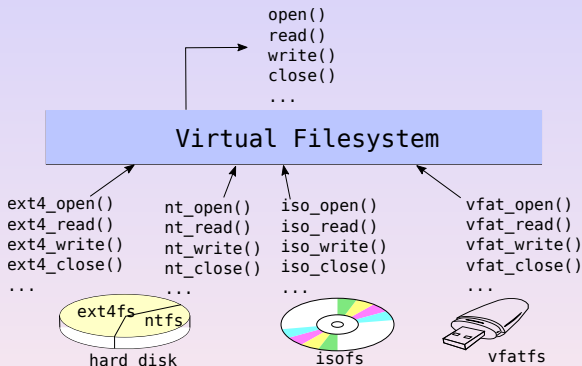
Evolution of X86

- 8086 (1979)
- 80286 (1982), Protected virtual address mode, multi-task
- 80386 (1985), 32bit(80386DX) protected mode, v86 mode supported.
- 80486 (1989), Internal cache(8KB) and math coprocessor(486DX), PLL.
- Pentium(1993), 8KB D-cache and 8KB P-cache, Superscaler pipelining, : SMM
- Pentium-MMX. Multi-Media eXtension : (or maybe Matrix Math eXtension)
- PentiumPro (1995), L1-cache and L2-cache.
- IA-64(Itanium, new ISA) vs. AMD64



VFS

Filesystem describes the way the data is stored and retrieved by filename.



VFS is a common interface for all filesystems.



Filesystems in Linux


Types of filesystem (Linux)

- Disk-based: FAT-series, EXT-series, YAFFS2, NTFS, ISO, ...
- Network: NFS, SAMBA, ...
- Pseudo filesystems: PROCFS, SYSFS, ...



Why Use FS

Non-FS structure




FILE_A FILE_B FILE_C FILE_D FILE_E
unused



Why Use FS

Filesystem

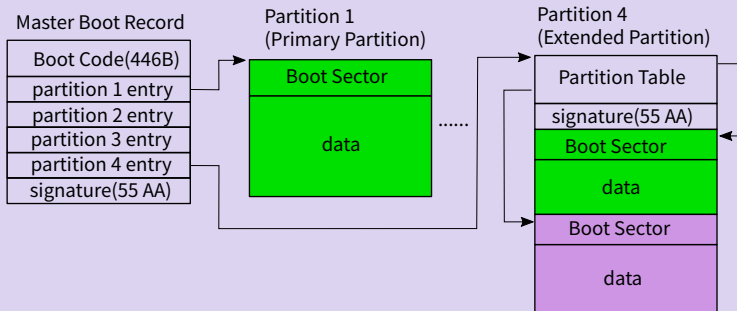


FILE_A
 FILE_B
 FILE_C
 FILE_D
 FILE_E
 unused



Disk Partition

MBR

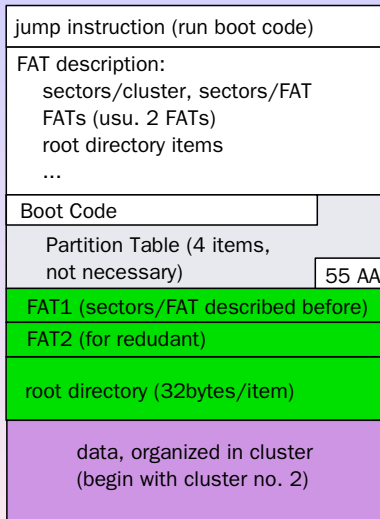


MBR uses 4 bytes describing a disk, max. 2^{41} bytes of a disk.

Modern disk uses GPT format, which uses LBA (Logical Block Address, 8 bytes), allows $(2^{64} - 1) \times 512 = 8\text{ZiB}$ of a disk size.



FAT File System



Partition Entry

80				Boot indicator
20	21	20		CHS (Start)
07				Filesystem type ID
59	1a	bf		CHS (End)
00	80	00	00	Relative sectors
00	0e	2e	00	Total sectors



FAT info

FAT format is described in boot sector:

offset	bytes	meaning
0000	3	JMP instruction
0003	8	Version info
000B	2	Bytes per sector
000D	1	Sectors per cluster
000E	2	Reserved sectors, incl.MBR
0010	1	Number of FATs
0011	2	Root entries, 32B/entry
0013	2	Sectors (on small volumes)
0015	1	Media Descriptor
0016	2	Sectors per FAT

Data begin at (Reserved sectors + FATs + Root entries), marked as cluster No.2.



FAT info(cont.)

offset	bytes	meaning
0018	2	Sectors per track
001A	2	Heads
001C	4	Hidden sectors
0020	4	Sectors (on large volumes)
0024	1	Physical driver number
0025	1	Reserved
0026	1	Signature (0x28 or 0x29)
0027	4	Serial number (random)
002A	11	Volume label ('NO NAME ')
0036	8	FS type ('FAT12 ')




Directory Entry

offset	bytes	meaning
00	11	Filename(8.3)
0B	1	Attribute
0C	10	Reserved
16	2	Modified time
18	2	Modified date
1A	2	First cluster
1C	4	File length(bytes)

Cluster mark: 000:free, 001:reserved,
002–FEF:next cluster number, FF0–FF6:
(reserved values), FF7: bad cluster, FF8–FFF:
last cluster
(Cluster No. 0 and 1 are for special purpose).



Ext2FS




Ext2 Layout

Inodes and block size are fixed once a disk is formatted. File attributes such as **size**, **permission**, **owner/group**, **location**, **date/time**, ... (but not filename) are stored in an inode.



INODE



Direct: $12 \times 1\text{KB}$

single indirect:
 $1\text{KB}/4 \times 1\text{KB} = 256\text{KB}$


doubly indirect:
 $256^2 \times 1\text{KB} = 64\text{MB}$

trebly indirect:
 $256^3 \times 1\text{KB} = 16\text{GB}$



Journal File System

System crash may lead to the loss of data when file written. File system can be corrupted if not closed when the system shuts down.
Using a journal allows data recovery of files and the data within it.



Summary

- Memory (RAM, ROM)
- Memory Organization
- Cache (SRAM vs. DRAM)
- Virtual Memory (Internal memory vs. Disk)
- Filesystem

